



#**RANKED 52**  
IN INDIA

#University Category



**NO.1** PVT. UNIVERSITY IN  
ACADEMIC REPUTATION IN INDIA



ACCREDITED **GRADE 'A'**  
BY NAAC



PERFECT SCORE OF **150/150** AS A TESTAMENT  
TO EXCEPTIONAL E-LEARNING METHODS



# Unit 3 : Inter Process Communication and Synchronization

**Dr. Abhijit and Mr. Deepak**

School of Computer Science  
UPES, Dehradun  
India

# Unit 3: Inter-Process Communication

L1	Inter Process Communication (IPC), IPC mechanisms: Shared Memory and Message Passing
L2	Message Passing: Shared Memory, Pipes and Named pipes in Linux
L3	Critical Section Problem, Race Condition,
L4	Producer Consumer Problem, Solution to Critical section Problem: Hardware and Software Solutions,
L5	Software Solutions: Semaphores: Counting semaphore, Binary semaphore,
L6	Software Solutions: Monitors
L7	Software Solutions: Algorithm 1, Algorithm 2,
L8	Software Solutions: Algorithm 3/Peterson Solution, Bakery Algorithm
L9	Classic process synchronization problems (case studies).

# Unit 3: Inter-Process Communication

## Lecture 1 | Inter Process Communication (IPC), IPC mechanisms: Shared Memory and Message Passing

### **Objective:**

- To discuss the communication mechanism among processes.
- To explore different approach to achieve inter process communication

# Interprocess Communication

- Interprocess Communication (IPC) refers to the mechanisms provided by the operating system that allow processes to communicate with each other and synchronize their actions.
- IPC is essential for the development of complex software systems where multiple processes need to work together and share data.
- Here are some common IPC mechanisms:
  - i) Pipes
  - ii) Message Passing
  - iii) Shared Memory
  - iv) Socket
  - v) Semaphores
  - vi) Remote Procedure Call (RPC)

# Interprocess Communication

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
  - Shared memory
  - Message passing

# Communications Models

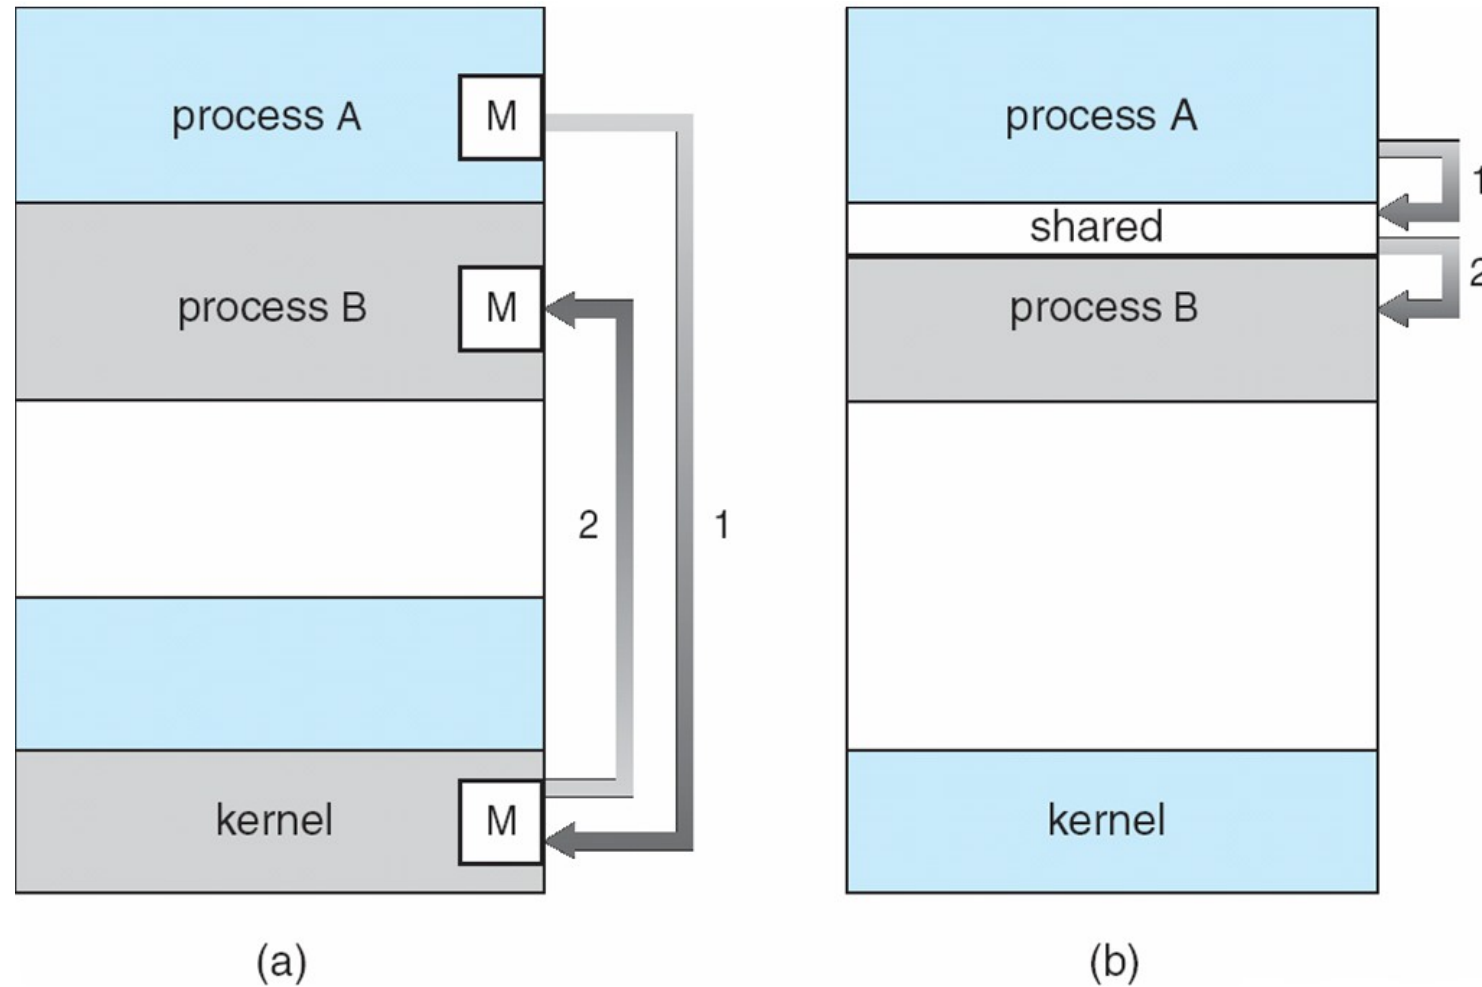


Fig 3.1 Interprocess communication model

# Cooperating Processes

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - Information sharing
  - Computation speed-up
  - Modularity
  - Convenience



# Examples of Cooperating Processes

- Web Servers: Multiple processes handle different client requests, sharing resources like memory and network connections.
- Database Systems: Processes collaborate to manage transactions, maintain data consistency, and ensure concurrency control.
- Distributed Systems: Processes running on different machines communicate and cooperate to perform distributed computations, share resources, and synchronize actions.
- Operating Systems: System processes cooperate to manage hardware resources, perform I/O operations, and execute user programs

# Benefits of Cooperating Processes

- Cooperating processes offer numerous benefits, including resource sharing, computation speedup, modularity, information sharing, improved reliability, enhanced communication, flexibility, and scalability.
- These advantages make cooperating processes essential in modern computing environments, particularly in multi-core, distributed, and real-time systems.

## Lecture 2 | Message Passing: Shared Memory, Pipes and Named pipes in Linux

### Objective:

- To understand the concepts of Interprocess Communication through Message passing
- To understand the concepts of Interprocess Communication through Shared memory
- To understand the concepts of Interprocess Communication through Pipes

# Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send**(*message*) – message size fixed or variable
  - **receive**(*message*)
- If  $P$  and  $Q$  wish to communicate, they need to:
  - establish a *communication link* between them
  - exchange messages via send/receive
- Implementation of communication link
  - physical (e.g., shared memory, hardware bus)
  - logical (e.g., logical properties)

# Interprocess Communication – Message Passing

- Examples of Message Passing in Different Contexts

1. Message Passing Interface (MPI) in Parallel Computing MPI is a standardized and portable message-passing system designed for parallel computing.

## An example using MPI in C

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    MPI_Init(&argc, &argv); // Initialize MPI environment
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank); // Get rank of the current process
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size); // Get total number of processes
    int number;
    if (world_rank == 0) {
        // Process 0 initializes the number
        number = -1;
        // Process 0 sends the number to process 1
        MPI_Send(&number, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
    } else if (world_rank == 1) {
        // Process 1 receives the number from process 0
        MPI_Recv(&number, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
        // Process 1 prints the received number
        printf("Process 1 received number %d from process 0\n", number);
    }
    MPI_Finalize(); // Finalize MPI environment
    return 0; }
```

# Pipes

- Acts as a conduit allowing two processes to communicate
- **Issues**
  - Is communication unidirectional or bidirectional?
  - In the case of two-way communication, is it half or full-duplex?
  - Must there exist a relationship (i.e. parent-child) between the communicating processes?
  - Can the pipes be used over a network?

# Ordinary Pipes

- **Ordinary Pipes** allow communication in standard producer-consumer style
- Producer writes to one end (the *write-end* of the pipe)
- Consumer reads from the other end (the *read-end* of the pipe)
- Ordinary pipes are therefore unidirectional
- Require parent-child relationship between communicating processes



# Ordinary Pipes

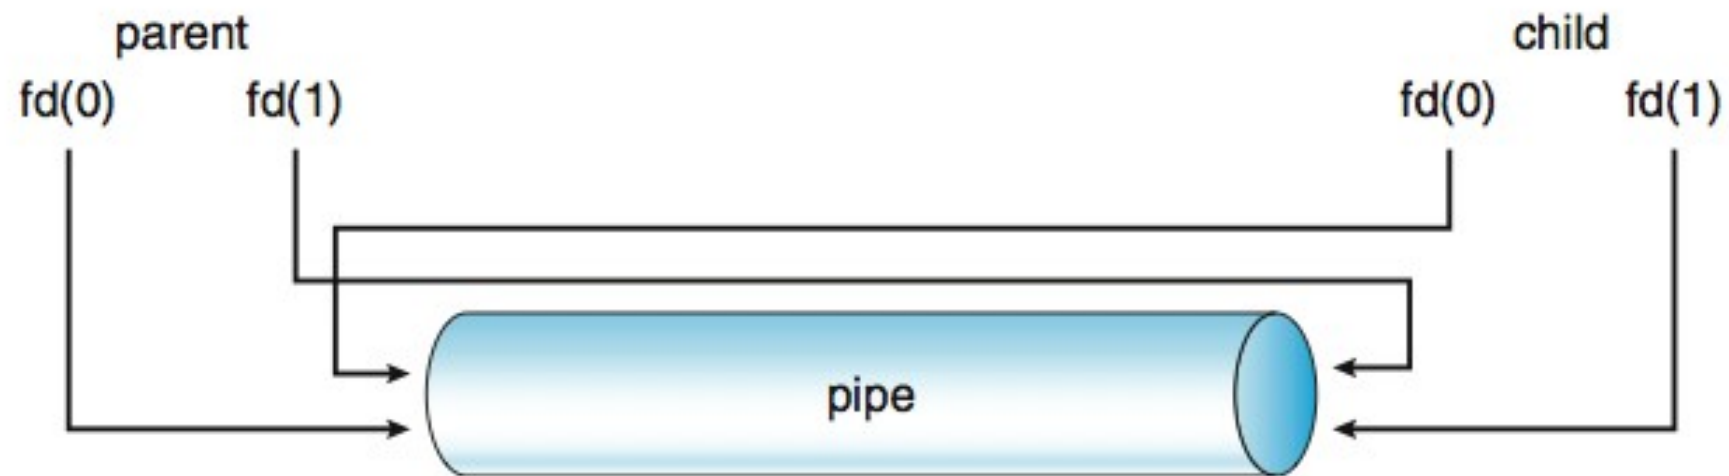


Fig: Ordinary Pipes

# Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems



# IPC mechanism in Pipes

```
#include <stdio.h>
#include <unistd.h>
int main() {
    int fd[2];
    pid_t pid;
    char buffer[30];
    pipe(fd);
    pid = fork();

    if (pid == 0) { // Child process
        close(fd[0]); // Close unused read end
        write(fd[1], "Hello, parent!", 14);
        close(fd[1]); // Close write end
    } else { // Parent process
        close(fd[1]); // Close unused write end
        read(fd[0], buffer, 14);
        printf("%s\n", buffer);
        close(fd[0]); // Close read end
    }
    return 0;
}
```

Objectives:

- To understand the concepts of Critical Section problem
- To understand the race condition

# Introduction

- Concurrent access to shared data may result in data inconsistency
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- Example: The producer-consumer problem.
  - Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **count** that keeps track of the number of full buffers. Initially, count is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

## Race Condition:

- A race condition occurs when multiple processes access and manipulate shared data concurrently, leading to unexpected results.
- Example: Incrementing and decrementing a shared counter.

# Race Condition

- `count++` could be implemented as

```
register1 = count  
register1 = register1 + 1  
count = register1
```

- `count--` could be implemented as

```
register2 = count  
register2 = register2 - 1  
count = register2
```

- Consider this execution interleaving with “count = 5” initially:

```
S0: producer execute register1 = count {register1 = 5}  
S1: producer execute register1 = register1 + 1 {register1 = 6}  
S2: consumer execute register2 = count {register2 = 5}  
S3: consumer execute register2 = register2 - 1 {register2 = 4}  
S4: producer execute count = register1 {count = 6}  
S5: consumer execute count = register2 {count = 4}
```

# Producer

```
while (true) {  
    /* produce an item and put in nextProduced */  
    while (count == BUFFER_SIZE)  
        ; // do nothing  
    buffer [in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;  
    count++;  
}
```

# Consumer

```
while (true) {  
    while (count == 0)  
        ; // do nothing  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    count--;  
  
    /* consume the item in  
    nextConsumed  
}
```



# Solution to Critical-Section Problem

1. Mutual Exclusion - If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections
2. Progress - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. Bounded Waiting - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
  - Assume that each process executes at a nonzero speed
  - No assumption concerning relative speed of the  $N$  processes

L4

## Producer Consumer Problem, Solution to Critical section Problem: Hardware and Software Solutions,

### **Objectives:**

- To understand the classical problems of Critical Section
- To device and develop the solution related to classical problems of Critical Section

# Peterson's Solution

- Two process solution
- Assume that the LOAD and STORE instructions are atomic; that is, cannot be interrupted.
- The two processes share two variables:
  - int **turn**;
  - Boolean **flag[2]**
- The variable **turn** indicates whose turn it is to enter the critical section.
- The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i]** = true implies that process  $P_i$  is ready!

# Algorithm for Process $P_i$

```
do {  
    flag[i] = TRUE;  
    turn = j;  
    while (flag[j] && turn == j);  
        critical section  
    flag[i] = FALSE;  
        remainder section  
} while (TRUE);
```

# Synchronization Hardware

- Many systems provide hardware support for critical section code
- Uniprocessors – could disable interrupts
  - Currently running code would execute without preemption
  - Generally too inefficient on multiprocessor systems
    - Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
  - Atomic = non-interruptable
  - Either test memory word and set value
  - Or swap contents of two memory words

## Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```

# TestAndSet Instruction

```
boolean TestAndSet (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

# Solution using TestAndSet

- Shared boolean variable lock., initialized to false.
- Solution:

```
do {  
    while ( TestAndSet (&lock ))  
        ; // do nothing  
  
    // critical section  
  
    lock = FALSE;  
  
    // remainder section  
  
} while (TRUE);
```



# Swap Instruction

- Definition:

```
void Swap (boolean *a, boolean *b)
{
    boolean temp = *a;
    *a = *b;
    *b = temp;
}
```

# Solution using Swap

- Shared Boolean variable lock initialized to FALSE; Each process has a local Boolean variable key
- Solution:

```
do {  
    key = TRUE;  
    while ( key == TRUE)  
        Swap (&lock, &key );  
  
    // critical section  
  
    lock = FALSE;  
  
    // remainder section  
  
} while (TRUE);
```

# Bounded-waiting Mutual Exclusion with TestandSet()

```
do {  
    waiting[i] = TRUE;  
    key = TRUE;  
    while (waiting[i] && key)  
        key = TestAndSet(&lock);  
    waiting[i] = FALSE;  
    // critical section  
    j = (i + 1) % n;  
    while ((j != i) && !waiting[j])  
        j = (j + 1) % n;  
    if (j == i)  
        lock = FALSE;  
    else  
        waiting[j] = FALSE;  
    // remainder section  
} while (TRUE);
```

L5

## Software Solutions: Semaphores: Counting semaphore, Binary semaphore,

### **Objective**

- To understand the concept of semaphores and their types
- To implement the different types of Semaphores

# Semaphore

- Semaphores are a synchronization mechanism used to control access to a common resource in concurrent programming and operating systems.
- They are particularly useful for managing access to shared resources and ensuring that multiple processes or threads do not simultaneously access a critical section, leading to race conditions or inconsistent data.

# Basic Operations: Semaphores

Support two atomic operations:

- wait (P or down): Decreases the semaphore value. If the value is already 0, the process or thread is blocked until the semaphore value becomes positive.
- signal (V or up): Increases the semaphore value. If there are any processes or threads waiting for the semaphore, one of them is unblocked.

wait(S):

while  $S \leq 0$ :

// Busy wait (or block the process/thread)

$S = S - 1$

signal(S):

$S = S + 1$

# Semaphore as General Synchronization Tool

- **Counting** semaphore – integer value can range over an unrestricted domain
- **Binary** semaphore – integer value can range only between 0 and 1; can be simpler to implement
  - Also known as **mutex locks**
- Can implement a counting semaphore **S** as a binary semaphore
- Provides mutual exclusion

```
Semaphore mutex; // initialized to 1
do {
    wait (mutex);
    // Critical Section
    signal (mutex);
    // remainder section
} while (TRUE);
```

# Semaphore Implementation

- Must guarantee that no two processes can execute `wait ()` and `signal ()` on the same semaphore at the same time
- Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section.
  - Could now have `busy waiting` in critical section implementation
    - But implementation code is short
    - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution.



## Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue. Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list
- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue.
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue.

# Semaphore Implementation with no Busy waiting (Cont.)

- Implementation of wait:

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

- Implementation of signal:

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

# Use of Semaphore: Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let  $s$  and  $q$  be two semaphores initialized to 1

$P_0$	$P_1$
wait (S);	wait (Q);
wait (Q);	wait (S);
.	.
.	.
.	.
signal (S);	signal (Q);
signal (Q);	signal (S);

- **Starvation** – indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** - Scheduling problem when lower-priority process holds a lock needed by higher-priority process

L6

## Software Solutions: Monitors

**Objective:**

- To explore the software solutions using Monitors

# Monitors

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...

    procedure Pn (...) {.....}

    Initialization code ( ....) { ... }
    ...
}
}
```

# Schematic view of a Monitor





Fig: Schematic View of Monitor

# Condition Variables

- `condition x, y;`
- Two operations on a condition variable:
  - `x.wait ()` – a process that invokes the operation is suspended.
  - `x.signal ()` – resumes one of processes (if any) that invoked `x.wait ()`

# Monitor with Condition Variables





# Solution to Dining Philosophers using Monitor

monitor DP

```
{  
    enum { THINKING; HUNGRY, EATING) state [5] ;  
    condition self [5];  
  
    void pickup (int i) {  
        state[i] = HUNGRY;  
        test(i);  
        if (state[i] != EATING) self [i].wait;  
    }  
  
    void putdown (int i) {  
        state[i] = THINKING;  
        // test left and right neighbors  
        test((i + 4) % 5);  
        test((i + 1) % 5);  
    }  
}
```

## Solution to Dining Philosophers (cont)

```
void test (int i) {  
    if ( (state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}  
  
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}
```

## Solution to Dining Philosophers (cont)

- Each philosopher / invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup (i);`

EAT

`DiningPhilosophers.putdown (i);`

# Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next-count = 0;
```

- Each procedure  $F$  will be replaced by

```
wait(mutex);
...
    body of  $F$ ;
...
if (next_count > 0)
    signal(next)
else
    signal(mutex);
```

- Mutual exclusion within a monitor is ensured.

# Monitor Implementation

- For each condition variable  $x$ , we have:

```
semaphore x_sem; // (initially = 0)  
int x-count = 0;
```

- The operation  $x.wait$  can be implemented as:

```
x-count++;  
if (next_count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x_sem);  
x-count--;
```

# Monitor Implementation

- The operation `x.signal` can be implemented as:

```
if (x-count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```

# A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = FALSE;
    }
}
```

## L7 & 8 | Software Solutions: Peterson Solution, Bakery Algorithm

Objective:

- To understand the software solutions using various algorithms



# Peterson's Solution

- Two-process solution
- Assumptions:
  - LOAD and STORE instructions are atomic (cannot be interrupted).
- Shared variables:
  - int turn;
  - boolean flag[2];
- Explanation:
  - turn indicates whose turn it is to enter the critical section.
  - flag array indicates if a process is ready to enter the critical section.  $\text{flag}[i] = \text{true}$  implies that process  $P_i$  is ready.

## Algorithm for Process $P_i$

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
    // Critical section  
    flag[i] = false;  
    // Remainder section  
} while (true);
```

# Bakery Algorithm

- N-process solution
- Assumptions:
  - The integers are assumed to be unbounded.
- Shared variables:
  - boolean choosing[N];
  - int number[N];
- Explanation:
  - choosing[i] indicates if process  $P_i$  is in the process of choosing a number.
  - number[i] holds the number chosen by process  $P_i$ . A higher number indicates a higher priority.

# Bakery Algorithm Cont..

```
do {
    choosing[i] = true;
    number[i] = 1 + max(number[0], number[1], ..., number[N-1]);
    choosing[i] = false;
    for (j = 0; j < N; j++) {
        while (choosing[j]);
        while ((number[j] != 0) && ((number[j], j) < (number[i], i)));
    }
    // Critical section
    number[i] = 0;
    // Remainder section
} while (true);
```

# Pthreads Synchronization

- Pthreads API is OS-independent
- It provides:
  - mutex locks
  - condition variables
- Non-portable extensions include:
  - read-write locks
  - spin locks

# Serializability

- Consider two data items A and B
- Consider Transactions  $T_0$  and  $T_1$
- Execute  $T_0$ ,  $T_1$  atomically
- Execution sequence called **schedule**
- Atomically executed transaction order called **serial schedule**
- For N transactions, there are  $N!$  valid serial schedules

# Schedule 1: $T_0$ then $T_1$

$T_0$	$T_1$
read( $A$ )	
write( $A$ )	
read( $B$ )	
write( $B$ )	
	read( $A$ )
	write( $A$ )
	read( $B$ )
	write( $B$ )

# Nonserial Schedule

- **Nonserial schedule** allows overlapped execute
  - Resulting execution not necessarily incorrect
- Consider schedule  $S$ , operations  $O_i, O_j$ 
  - **Conflict** if access same data item, with at least one write
- If  $O_i, O_j$  consecutive and operations of different transactions &  $O_i$  and  $O_j$  don't conflict
  - Then  $S'$  with swapped order  $O_j O_i$  equivalent to  $S$
- If  $S$  can become  $S'$  via swapping nonconflicting operations
  - $S$  is **conflict serializable**

## Schedule 2: Concurrent Serializable Schedule

$T_0$	$T_1$
read( $A$ )	
write( $A$ )	
	read( $A$ )
	write( $A$ )
read( $B$ )	
write( $B$ )	
	read( $B$ )
	write( $B$ )



# Locking Protocol

- Ensure serializability by associating lock with each data item
  - Follow locking protocol for access control
- Locks
  - **Shared** –  $T_i$  has shared-mode lock (S) on item Q,  $T_i$  can read Q but not write Q
  - **Exclusive** –  $T_i$  has exclusive-mode lock (X) on Q,  $T_i$  can read and write Q
- Require every transaction on item Q acquire appropriate lock
- If lock already held, new request may have to wait
  - Similar to readers-writers algorithm

# Two-phase Locking Protocol

- Generally ensures conflict serializability
- Each transaction issues lock and unlock requests in two phases
  - Growing – obtaining locks
  - Shrinking – releasing locks
- Does not prevent deadlock

# Timestamp-based Protocols

- Select order among transactions in advance – **timestamp-ordering**
- Transaction  $T_i$  associated with timestamp  $TS(T_i)$  before  $T_i$  starts
  - $TS(T_i) < TS(T_j)$  if  $T_i$  entered system before  $T_j$
  - TS can be generated from system clock or as logical counter incremented at each entry of transaction
- Timestamps determine serializability order
  - If  $TS(T_i) < TS(T_j)$ , system must ensure produced schedule equivalent to serial schedule where  $T_i$  appears before  $T_j$

# Timestamp-based Protocol Implementation

- Data item Q gets two timestamps
  - W-timestamp(Q) – largest timestamp of any transaction that executed write(Q) successfully
  - R-timestamp(Q) – largest timestamp of successful read(Q)
  - Updated whenever read(Q) or write(Q) executed
- **Timestamp-ordering protocol** assures any conflicting **read** and **write** executed in timestamp order
- Suppose  $T_i$  executes **read(Q)**
  - If  $TS(T_i) < W\text{-timestamp}(Q)$ ,  $T_i$  needs to read value of Q that was already overwritten
    - **read** operation rejected and  $T_i$  rolled back
  - If  $TS(T_i) \geq W\text{-timestamp}(Q)$ 
    - **read** executed, R-timestamp(Q) set to  $\max(R\text{-timestamp}(Q), TS(T_i))$

# Timestamp-ordering Protocol

- Suppose  $T_i$  executes `write(Q)`
  - If  $TS(T_i) < R\text{-timestamp}(Q)$ , value  $Q$  produced by  $T_i$  was needed previously and  $T_i$  assumed it would never be produced
    - `Write` operation rejected,  $T_i$  rolled back
  - If  $TS(T_i) < W\text{-timestamp}(Q)$ ,  $T_i$  attempting to write obsolete value of  $Q$ 
    - `Write` operation rejected and  $T_i$  rolled back
  - Otherwise, `write` executed
- Any rolled back transaction  $T_i$  is assigned new timestamp and restarted
- Algorithm ensures conflict serializability and freedom from deadlock

## Schedule Possible Under Timestamp Protocol

$T_2$	$T_3$
read( $B$ )	
	read( $B$ )
	write( $B$ )
read( $A$ )	
	read( $A$ )
	write( $A$ )

L9

## Classic process synchronization problems (case studies)

Objective:

- To understand the classic Synchronization examples and case studies

# Synchronization Examples

- Solaris
- Windows XP
- Linux
- Pthreads



# Solaris Synchronization

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing
- Uses **adaptive mutexes** for efficiency when protecting data from short code segments
- Uses **condition variables** and **readers-writers** locks when longer sections of code need access to data
- Uses **turnstile** to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock

# Windows XP Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems
- Uses **spinlocks** on multiprocessor systems
- Also provides **dispatcher objects** which may act as either mutexes and semaphores
- Dispatcher objects may also provide **events**
  - An event acts much like a condition variable

# Linux Synchronization

- Linux:
  - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
  - Version 2.6 and later, fully preemptive
- Linux provides:
  - semaphores
  - spin locks

# Classical Problems of Synchronization

- Bounded-Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

# Bounded-Buffer Problem

- $N$  buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value  $N$ .

# Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do {  
  
    // produce an item in nextp  
  
    wait (empty);  
    wait (mutex);  
  
    // add the item to the buffer  
  
    signal (mutex);  
    signal (full);  
} while (TRUE);
```

# Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
do {  
    wait (full);  
    wait (mutex);  
  
    // remove an item from buffer to nextc  
  
    signal (mutex);  
    signal (empty);  
  
    // consume the item in nextc  
  
} while (TRUE);
```

# Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write
- Problem – allow multiple readers to read at the same time. Only one single writer can access the shared data at the same time
- Shared Data
  - Data set
  - Semaphore **mutex** initialized to 1
  - Semaphore **wrt** initialized to 1
  - Integer **readcount** initialized to 0



# Readers-Writers Problem (Cont.)

- The structure of a writer process


```
do {  
    wait (wrt) ;  
  
    //  writing is performed  
  
    signal (wrt) ;  
} while (TRUE);
```

# Readers-Writers Problem (Cont.)

- The structure of a reader process

```
do {  
    wait (mutex) ;  
    readcount ++ ;  
    if (readcount == 1)  
        wait (wrt) ;  
    signal (mutex)  
    // reading is performed  
    wait (mutex) ;  
    readcount -- ;  
    if (readcount == 0)  
        signal (wrt) ;  
    signal (mutex) ;  
} while (TRUE);
```

# Dining-Philosophers Problem



- Shared data
  - Bowl of rice (data set)
  - Semaphore **chopstick** [5] initialized to 1

# Dining-Philosophers Problem (Cont.)

- The structure of Philosopher  $i$ :

```
do {  
    wait ( chopstick[i] );  
    wait ( chopStick[ (i + 1) % 5] );  
  
    // eat  
  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (TRUE);
```

# Problems with Semaphores

- Incorrect use of semaphore operations:
  - signal (mutex) .... wait (mutex)
  - wait (mutex) ... wait (mutex)
  - Omitting of wait (mutex) or signal (mutex) (or both)

# Multiple-Choice Questions (MCQs)

1) Which of the following is NOT a mechanism for Inter-Process Communication (IPC)?

- A) Shared Memory
- B) Message Passing
- C) Pipelining
- D) Remote Procedure Calls (RPC)

2) In message passing, what is the function of the send(message) operation?

- A) To receive a message from another process
- B) To send a message to another process
- C) To create a shared memory space
- D) To initialize message buffers

3) What is a race condition?

- A) A condition where processes race to complete tasks
- B) A scenario where two processes compete for CPU time
- C) A condition where the outcome depends on the non-deterministic ordering of processes
- D) A situation where processes run in a round-robin manner

4) Which of the following is NOT a requirement for the critical-section problem solution?

- A) Mutual Exclusion
- B) Progress
- C) Bounded Waiting
- D) Unbounded Execution

5) What is the primary role of a semaphore in process synchronization?

- A) To allocate CPU time to processes
- B) To manage memory allocation
- C) To ensure mutual exclusion and synchronization
- D) To handle I/O operations

6) Which synchronization mechanism uses wait() and signal() operations?

- A) Locks
- B) Semaphores
- C) Monitors
- D) Barriers

# Multiple-Choice Questions (MCQs)

7) In Peterson's solution for the critical-section problem, what does the turn variable signify?

- A) The process currently in the critical section
- B) The process allowed to enter the critical section next
- C) The number of processes waiting to enter the critical section
- D) The priority of processes

8) What is a monitor in the context of process synchronization?

- A) A low-level synchronization primitive
- B) A high-level synchronization abstraction
- C) A type of semaphore
- D) A hardware-based synchronization mechanism

9) Which of the following problems is an example of a classic synchronization problem?

- A) Memory Allocation Problem
- B) Process Scheduling Problem
- C) Dining Philosophers Problem
- D) Disk Scheduling Problem

10) Which of the following operating systems uses adaptive mutexes for synchronization?

- A) Windows XP
- B) Solaris
- C) Linux
- D) MacOS

## Answers:

1. C) Pipelining
2. B) To send a message to another process
3. C) A condition where the outcome depends on the non-deterministic ordering of processes
4. D) Unbounded Execution
5. C) To ensure mutual exclusion and synchronization
6. B) Semaphores
7. B) The process allowed to enter the critical section next
8. B) A high-level synchronization abstraction
9. C) Dining Philosophers Problem
10. B) Solaris





**Thank You**

