

# Lab File

## Fundamentals of Data Science Lab

**Session: Jan - May 2025**

Programme: BTech. CS - Data Science

Sem: 4

Batch: 5

**Submitted By:**

Name: Kshitij Chandrakar  
SAP ID: 500124827

**Submitted To:**

Dr. Sachi Chaudhary

# Experiment 6

## Excercise1

April 21, 2025

```
[1]: library(ggplot2)
      library(dplyr)
      library(cluster)
```

Attaching package: ‘dplyr’

The following objects are masked from ‘package:stats’:

filter, lag

The following objects are masked from ‘package:base’:

intersect, setdiff, setequal, union

```
[2]: set.seed(1)
```

```
[3]: mall_data <- read.csv("/home/asus/content/Notes/Semester 4/FDN Lab/Experiments/
      ↪Experiment 6/Mall_Customers.csv")
```

```
[4]: head(mall_data)
      summary(mall_data)
```

|                     |        | CustomerID | Gender | Age    | Annual.Income..k.. | Spending.Score..1.100. |
|---------------------|--------|------------|--------|--------|--------------------|------------------------|
|                     |        | <int>      | <chr>  | <int>  | <int>              | <int>                  |
| A data.frame: 6 × 5 | 1      | 1          | Male   | 19     | 15                 | 39                     |
|                     | 2      | 2          | Male   | 21     | 15                 | 81                     |
|                     | 3      | 3          | Female | 20     | 16                 | 6                      |
|                     | 4      | 4          | Female | 23     | 16                 | 77                     |
|                     | 5      | 5          | Female | 31     | 17                 | 40                     |
|                     | 6      | 6          | Female | 22     | 17                 | 76                     |
| CustomerID          |        | Gender     | Age    |        | Annual.Income..k.. |                        |
| Min.                | : 1.00 | Length:200 | Min.   | :18.00 | Min.               | : 15.00                |

|                |                  |               |                |
|----------------|------------------|---------------|----------------|
| 1st Qu.: 50.75 | Class :character | 1st Qu.:28.75 | 1st Qu.: 41.50 |
| Median :100.50 | Mode :character  | Median :36.00 | Median : 61.50 |
| Mean :100.50   |                  | Mean :38.85   | Mean : 60.56   |
| 3rd Qu.:150.25 |                  | 3rd Qu.:49.00 | 3rd Qu.: 78.00 |
| Max. :200.00   |                  | Max. :70.00   | Max. :137.00   |

Spending.Score..1.100.

Min. : 1.00

1st Qu.:34.75

Median :50.00

Mean :50.20

3rd Qu.:73.00


Max. :99.00

```
[5]: colnames(mall_data) <- c("CustomerID", "Gender", "Age", "AnnualIncome", "SpendingScore")
```

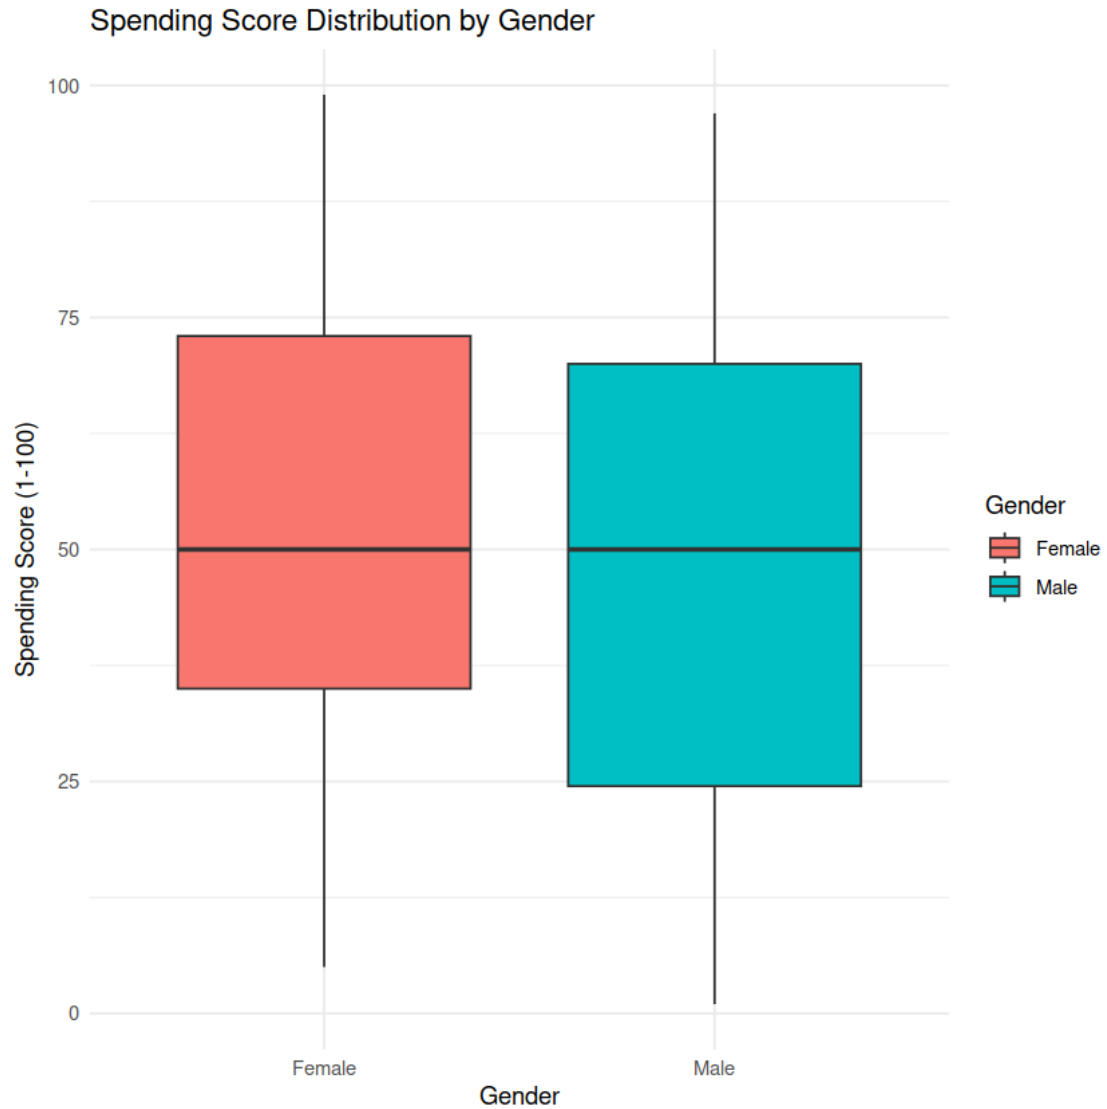
```
[6]: ggplot(mall_data, aes(x = AnnualIncome, y = SpendingScore)) +
  geom_point(aes(color = Gender), alpha = 0.7) +
  labs(title = "Annual Income vs Spending Score by Gender",
       x = "Annual Income (k$)",
       y = "Spending Score (1-100)") +
  theme_minimal()
```



```
[7]: ggplot(mall_data, aes(x = Age)) +
  geom_histogram(binwidth = 5, fill = "skyblue", color = "black") +
  labs(title = "Age Distribution of Customers",
    x = "Age",
    y = "Count") +
  theme_minimal()
```



```
[8]: ggplot(mall_data, aes(x = Gender, y = SpendingScore, fill = Gender)) +
  geom_boxplot() +
  labs(title = "Spending Score Distribution by Gender",
        y = "Spending Score (1-100)") +
  theme_minimal()
# =====
# Part 2
# =====
```




```
[9]: data_for_clustering <- mall_data[, c("AnnualIncome", "SpendingScore")]
```

```
[10]: wss <- function(k) {
      kmeans(data_for_clustering, k, nstart = 10)$tot.withinss
    }
```

```
[11]: k_values <- 1:10
      wss_values <- sapply(k_values, wss)
```

```
[12]: plot(k_values, wss_values,
          type = "b", pch = 19, frame = FALSE,
          xlab = "Number of clusters K",
          ylab = "Total within-clusters sum of squares",
```

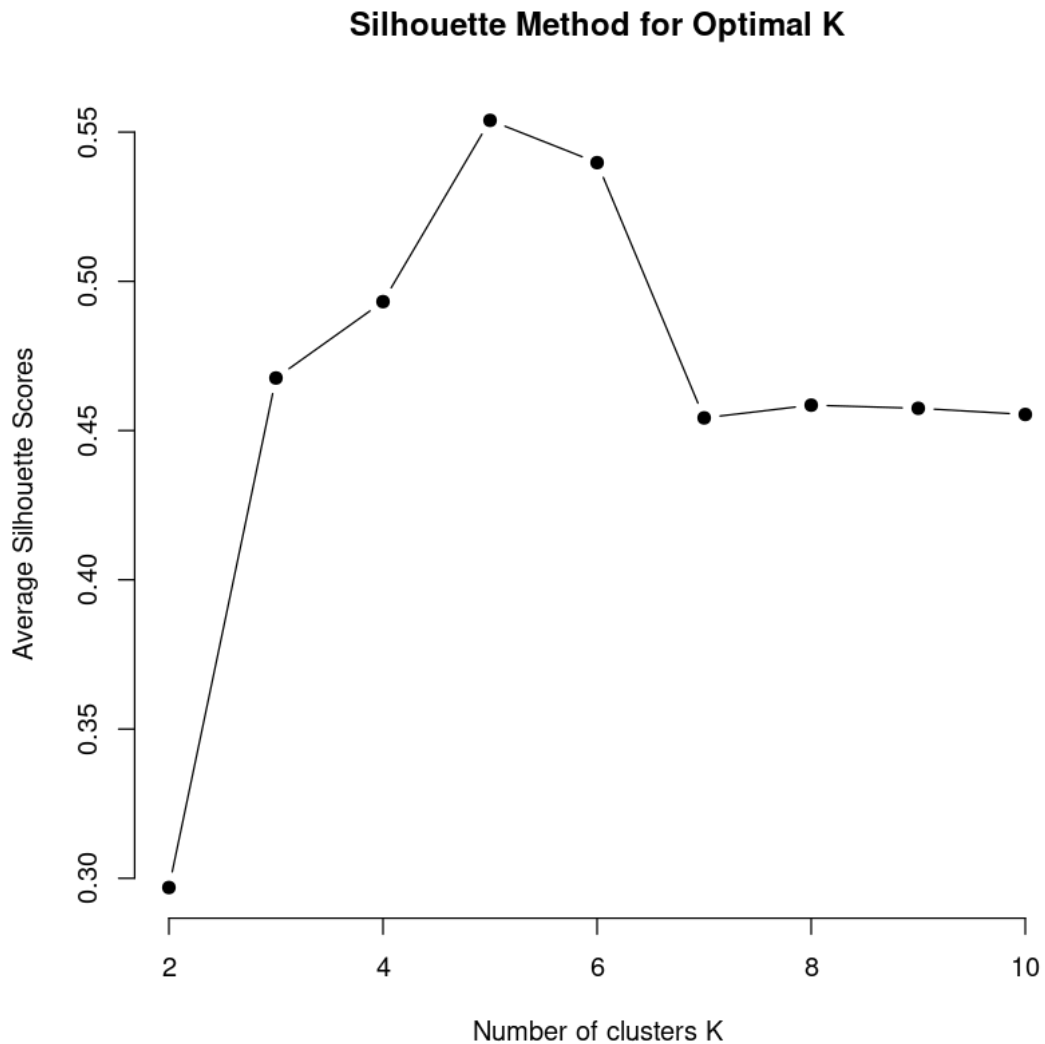
```
main = "Elbow Method for Optimal K")
```



```
[13]: avg_sil <- function(k) {  
  km.res <- kmeans(data_for_clustering, centers = k, nstart = 25)  
  ss <- silhouette(km.res$cluster, dist(data_for_clustering))  
  mean(ss[, 3])  
}
```

```
[14]: k_values <- 2:10  
avg_sil_values <- sapply(k_values, avg_sil)
```

```
[15]: plot(k_values, avg_sil_values,
          type = "b", pch = 19, frame = FALSE,
          xlab = "Number of clusters K",
          ylab = "Average Silhouette Scores",
          main = "Silhouette Method for Optimal K")
```




```
[16]: final_k <- 5
```

```
[17]: kmeans_result <- kmeans(data_for_clustering, centers = final_k, nstart = 25)
mall_data$Cluster <- as.factor(kmeans_result$cluster)
```



```
[18]: ggplot(mall_data, aes(x = AnnualIncome, y = SpendingScore, color = Cluster)) +
  geom_point(size = 3, alpha = 0.7) +
  scale_color_brewer(palette = "Set1") +
  labs(title = "Customer Segments by Annual Income and Spending Score",
       x = "Annual Income (k$)",
       y = "Spending Score (1-100)") +
  theme_minimal()
```



```
[19]: cluster_stats <- mall_data %>%
  group_by(Cluster) %>%
  summarise(
    Count = n(),
    Avg_Age = mean(Age),
```

8

```

    Avg_Income = mean(AnnualIncome),
    Avg_Spending = mean(SpendingScore),
    Female_Pct = sum(Gender == "Female") / n() * 100
  )

```

```

[20]: # Printing Stats
      print(cluster_stats)

```

```

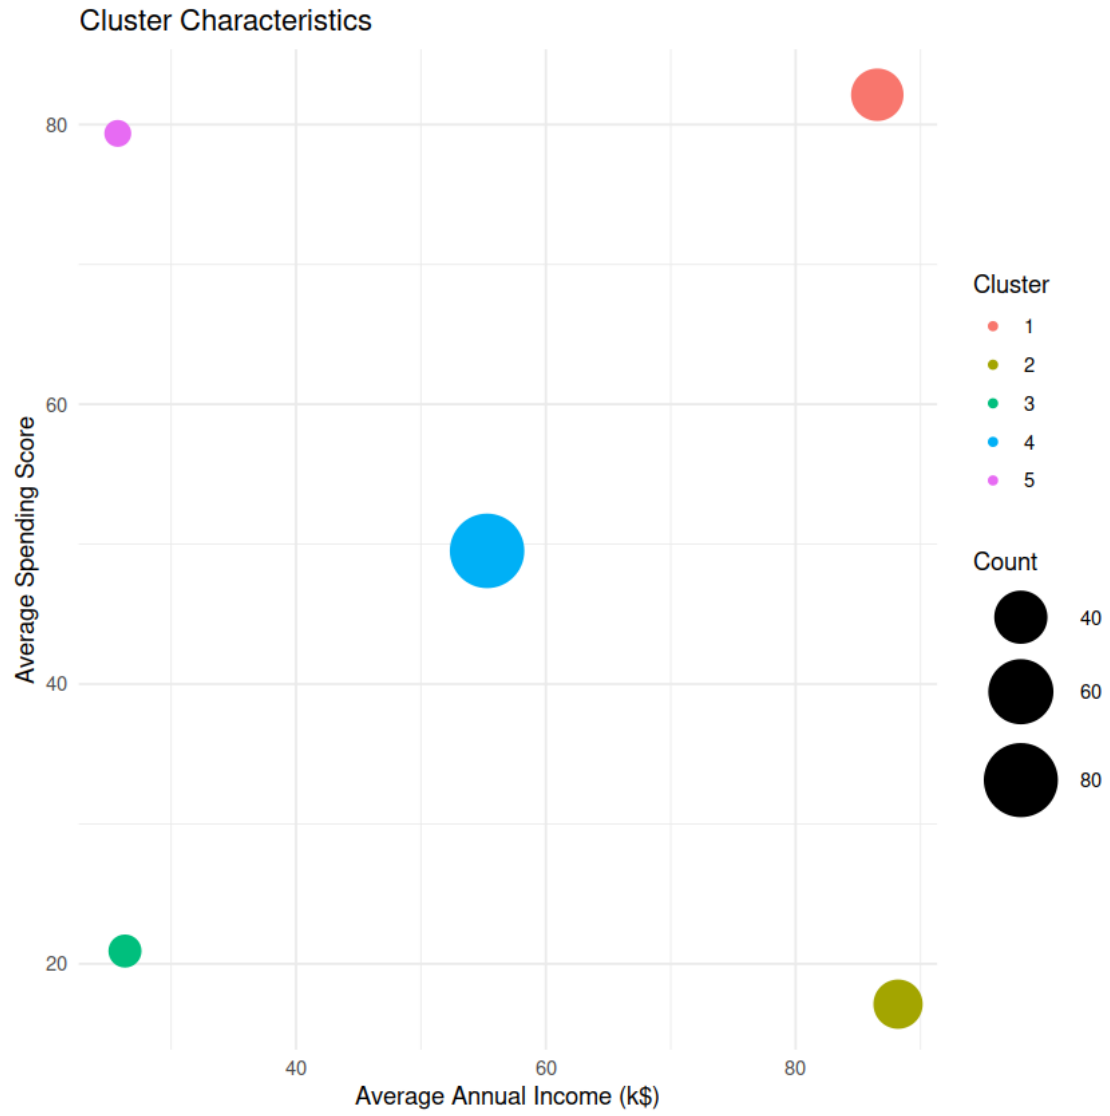
# A tibble: 5 × 6
  Cluster Count Avg_Age Avg_Income Avg_Spending Female_Pct
  <dbl>    <dbl>   <dbl>    <dbl>      <dbl>
1 1         39    32.7     86.5      82.1     53.8
2 2         35    41.1     88.2      17.1     45.7
3 3         23    45.2     26.3      20.9     60.9
4 4         81    42.7     55.3      49.5     59.3
5 5         22    25.3     25.7      79.4     59.1

```

```

[21]: ggplot(cluster_stats, aes(x = Avg_Income, y = Avg_Spending, size = Count, color_
      ↪= Cluster)) +
      geom_point() +
      scale_size(range = c(5, 15)) +
      labs(title = "Cluster Characteristics",
           x = "Average Annual Income (k$)",
           y = "Average Spending Score") +
      theme_minimal()

```



## Excercise2

April 21, 2025

```
[8]: library(tidyverse)
      library(cluster)
```

```
[9]: set.seed(1)
```

```
[10]: happiness <- read_csv("/home/asus/content/Notes/Semester 4/FDN Lab/Experiments/
      ↪Experiment 6/archive(11)/2019.csv")
```

Rows: 156 Columns: 9  
Column specification

Delimiter: ","

chr (1): Country

dbl (8): Overall rank, Score, GDP per capita, Social support, Healthy  
life e...

Use `spec()` to retrieve the full column specification for this  
data.

Specify the column types or set `show\_col\_types = FALSE` to quiet  
this message.

```
[11]: head(happiness)
      summary(happiness)
```

|                 | Overall rank | Country     | Score | GDP per capita | Social support | Healthy life expectancy |
|-----------------|--------------|-------------|-------|----------------|----------------|-------------------------|
|                 | <dbl>        | <chr>       | <dbl> | <dbl>          | <dbl>          | <dbl>                   |
| A tibble: 6 × 9 | 1            | Finland     | 7.769 | 1.340          | 1.587          | 0.986                   |
|                 | 2            | Denmark     | 7.600 | 1.383          | 1.573          | 0.996                   |
|                 | 3            | Norway      | 7.554 | 1.488          | 1.582          | 1.028                   |
|                 | 4            | Iceland     | 7.494 | 1.380          | 1.624          | 1.026                   |
|                 | 5            | Netherlands | 7.488 | 1.396          | 1.522          | 0.999                   |
|                 | 6            | Switzerland | 7.480 | 1.452          | 1.526          | 1.052                   |

| Overall rank   | Country          | Score          | GDP per capita  |
|----------------|------------------|----------------|-----------------|
| Min. : 1.00    | Length:156       | Min. : 2.853   | Min. : 0.0000   |
| 1st Qu.: 39.75 | Class :character | 1st Qu.: 4.545 | 1st Qu.: 0.6028 |
| Median : 78.50 | Mode :character  | Median : 5.380 | Median : 0.9600 |

|                |                           |                              |         |          |         |
|----------------|---------------------------|------------------------------|---------|----------|---------|
| Mean           | : 78.50                   | Mean                         | :5.407  | Mean     | :0.9051 |
| 3rd Qu.:       | 117.25                    | 3rd Qu.:                     | 6.184   | 3rd Qu.: | 1.2325  |
| Max.           | :156.00                   | Max.                         | :7.769  | Max.     | :1.6840 |
| Social support | Healthy life expectancy   | Freedom to make life choices |         |          |         |
| Min.           | :0.000                    | Min.                         | :0.0000 | Min.     | :0.0000 |
| 1st Qu.:       | 1.056                     | 1st Qu.:                     | 0.5477  | 1st Qu.: | 0.3080  |
| Median         | :1.272                    | Median                       | :0.7890 | Median   | :0.4170 |
| Mean           | :1.209                    | Mean                         | :0.7252 | Mean     | :0.3926 |
| 3rd Qu.:       | 1.452                     | 3rd Qu.:                     | 0.8818  | 3rd Qu.: | 0.5072  |
| Max.           | :1.624                    | Max.                         | :1.1410 | Max.     | :0.6310 |
| Generosity     | Perceptions of corruption |                              |         |          |         |
| Min.           | :0.0000                   | Min.                         | :0.0000 |          |         |
| 1st Qu.:       | 0.1087                    | 1st Qu.:                     | 0.0470  |          |         |
| Median         | :0.1775                   | Median                       | :0.0855 |          |         |
| Mean           | :0.1848                   | Mean                         | :0.1106 |          |         |
| 3rd Qu.:       | 0.2482                    | 3rd Qu.:                     | 0.1412  |          |         |
| Max.           | :0.5660                   | Max.                         | :0.4530 |          |         |

```
[12]: # Filtering out the text cols
features <- c("Overall rank", "Score", "GDP per capita", "Social support",
  ↪ "Healthy life expectancy", "Freedom to make life choices", "Generosity",
  ↪ "Perceptions of corruption")

happiness <- happiness %>%
  select(all_of(features)) %>%
  na.omit()
```

```
[13]: features <- c("Overall rank", "Score", "GDP per capita", "Social support",
  ↪ "Healthy life expectancy", "Freedom to make life choices", "Generosity",
  ↪ "Perceptions of corruption")

happiness_country <- happiness %>%
  select(all_of(features))
```

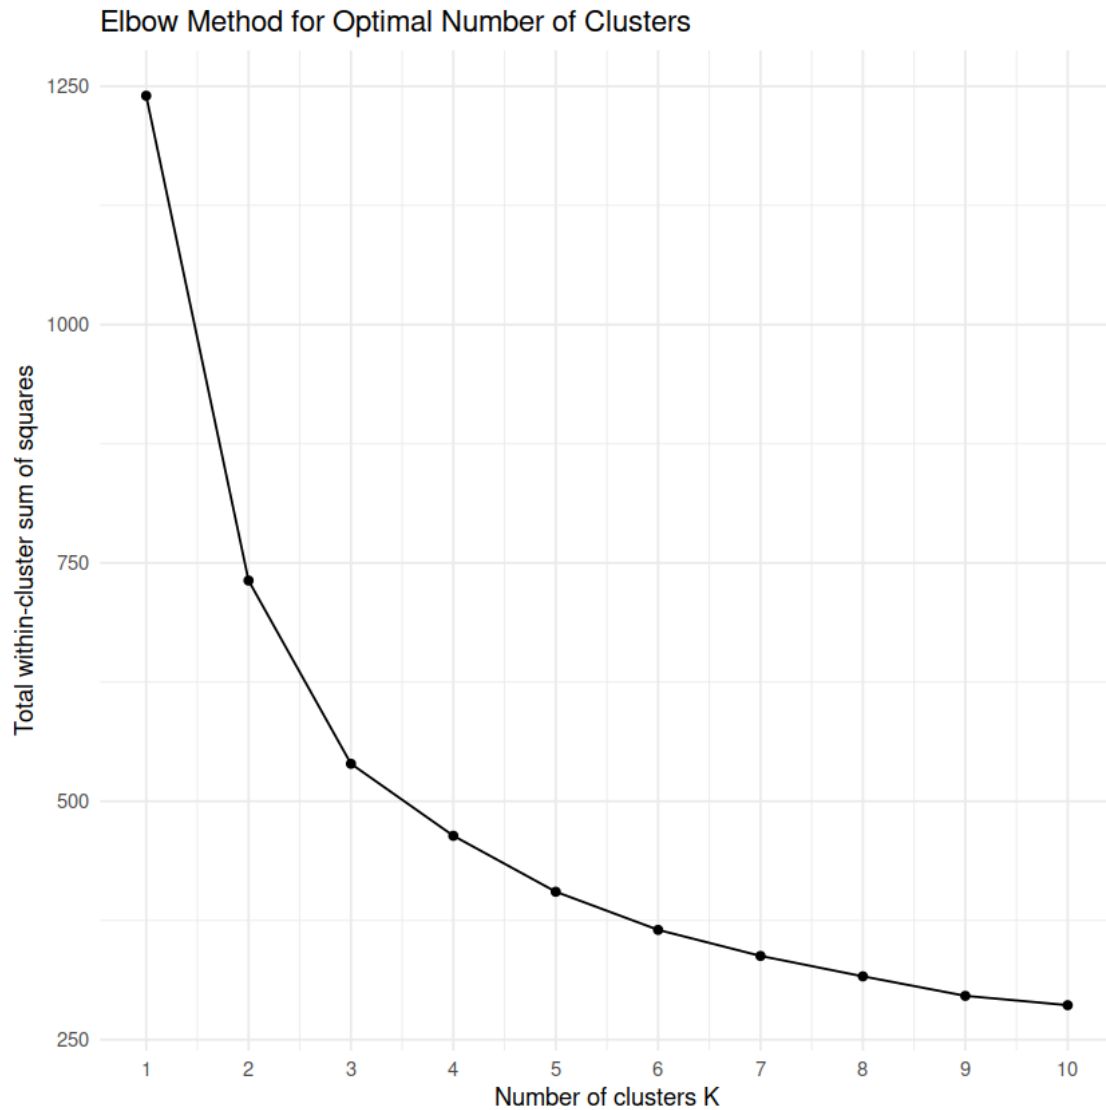
```
[14]: #SCAAAAAAAAAAAAAAAAALIng
happiness_scaled <- scale(happiness)
```

```
[15]: # total within-cluster sum of squares
wss <- function(k) {
  kmeans(happiness_scaled, k, nstart = 10)$tot.withinss
}
```

```
[16]: k_values <- 1:10
wss_values <- map_dbl(k_values, wss)
```

```
[17]: ggplot(data.frame(k = k_values, wss = wss_values), aes(k, wss)) +
  geom_line() + geom_point() +
  12
```


```
scale_x_continuous(breaks = k_values) +
labs(title = "Elbow Method for Optimal Number of Clusters",
      x = "Number of clusters K",
      y = "Total within-cluster sum of squares") +
theme_minimal()
```



```
[18]: avg_sil <- function(k) {
  km.res <- kmeans(happiness_scaled, centers = k, nstart = 25)
  ss <- silhouette(km.res$cluster, dist(happiness_scaled))
  mean(ss[, 3])
}
```

```
[19]: k_values <- 2:10
      avg_sil_values <- map_dbl(k_values, avg_sil)
```

```
[20]: ggplot(data.frame(k = k_values, silhouette = avg_sil_values), aes(k, silhouette)) +
      geom_line() + geom_point() +
      scale_x_continuous(breaks = k_values) +
      labs(title = "Silhouette Method for Optimal Number of Clusters",
           x = "Number of clusters K",
           y = "Average Silhouette Width") +
      theme_minimal()
```



```
[21]: optimal_k <- 5
```

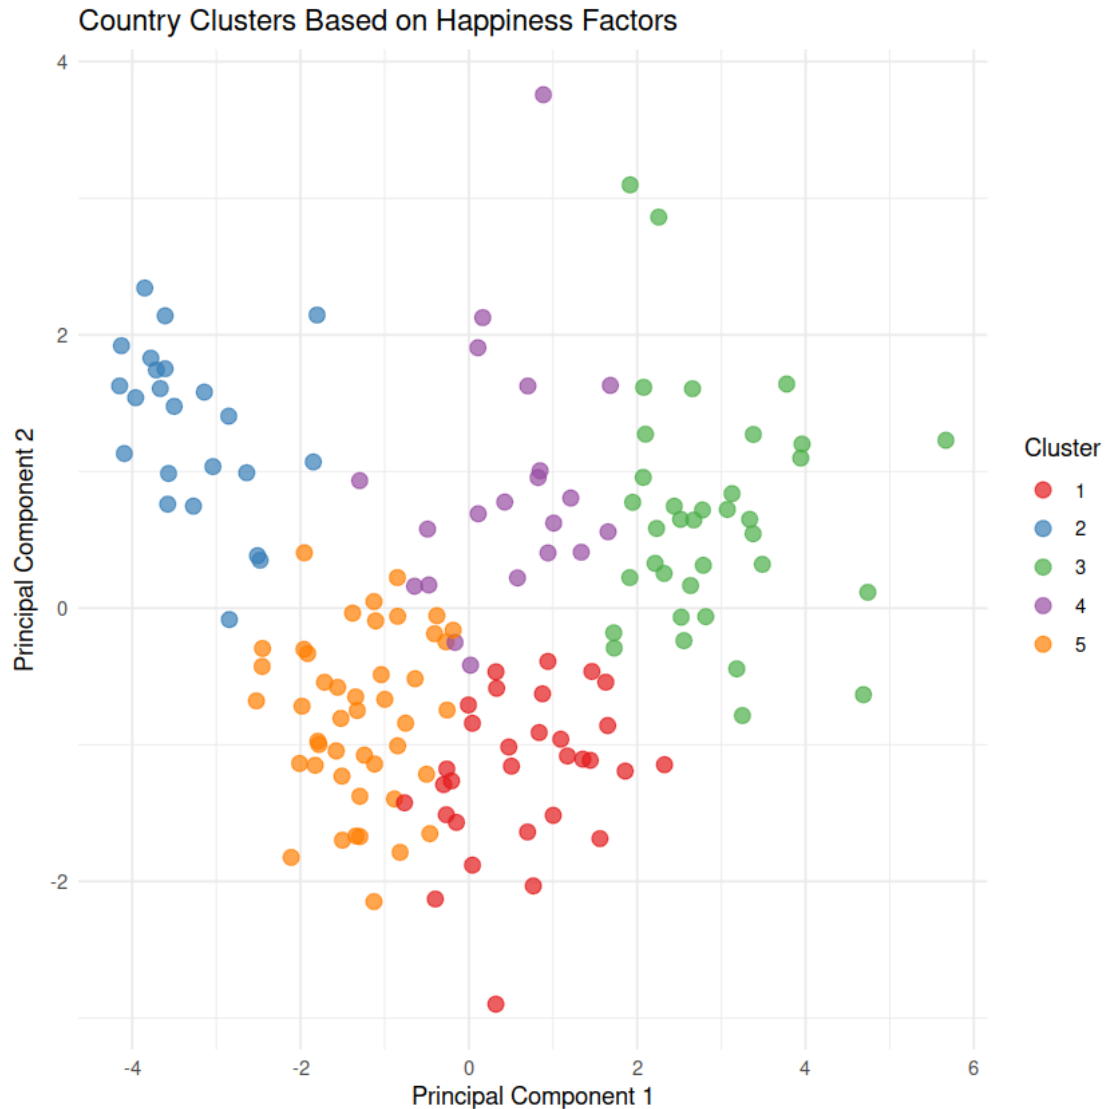
```
[22]: kmeans_result <- kmeans(happiness_scaled, centers = optimal_k, nstart = 25)
```

```
[23]: # Add cluster assignments to original data  
happiness$Cluster <- as.factor(kmeans_result$cluster)
```

```
[24]: # Perform PCA for visualization  
pca_result <- prcomp(happiness_scaled, scale. = TRUE)  
pca_df <- as.data.frame(pca_result$x[, 1:2])  
pca_df$Cluster <- happiness$Cluster
```

```
[25]: # Plot clusters in PCA space  
ggplot(pca_df, aes(x = PC1, y = PC2, color = Cluster)) +  
  geom_point(size = 3, alpha = 0.7) +  
  scale_color_brewer(palette = "Set1") +  
  labs(title = "Country Clusters Based on Happiness Factors",  
        x = "Principal Component 1",  
        y = "Principal Component 2") +  
  theme_minimal()
```





```
[26]: # Prepare data for parallel coordinates plot
cluster_means <- happiness %>%
  group_by(Cluster) %>%
  summarise(across(where(is.numeric), mean))
```

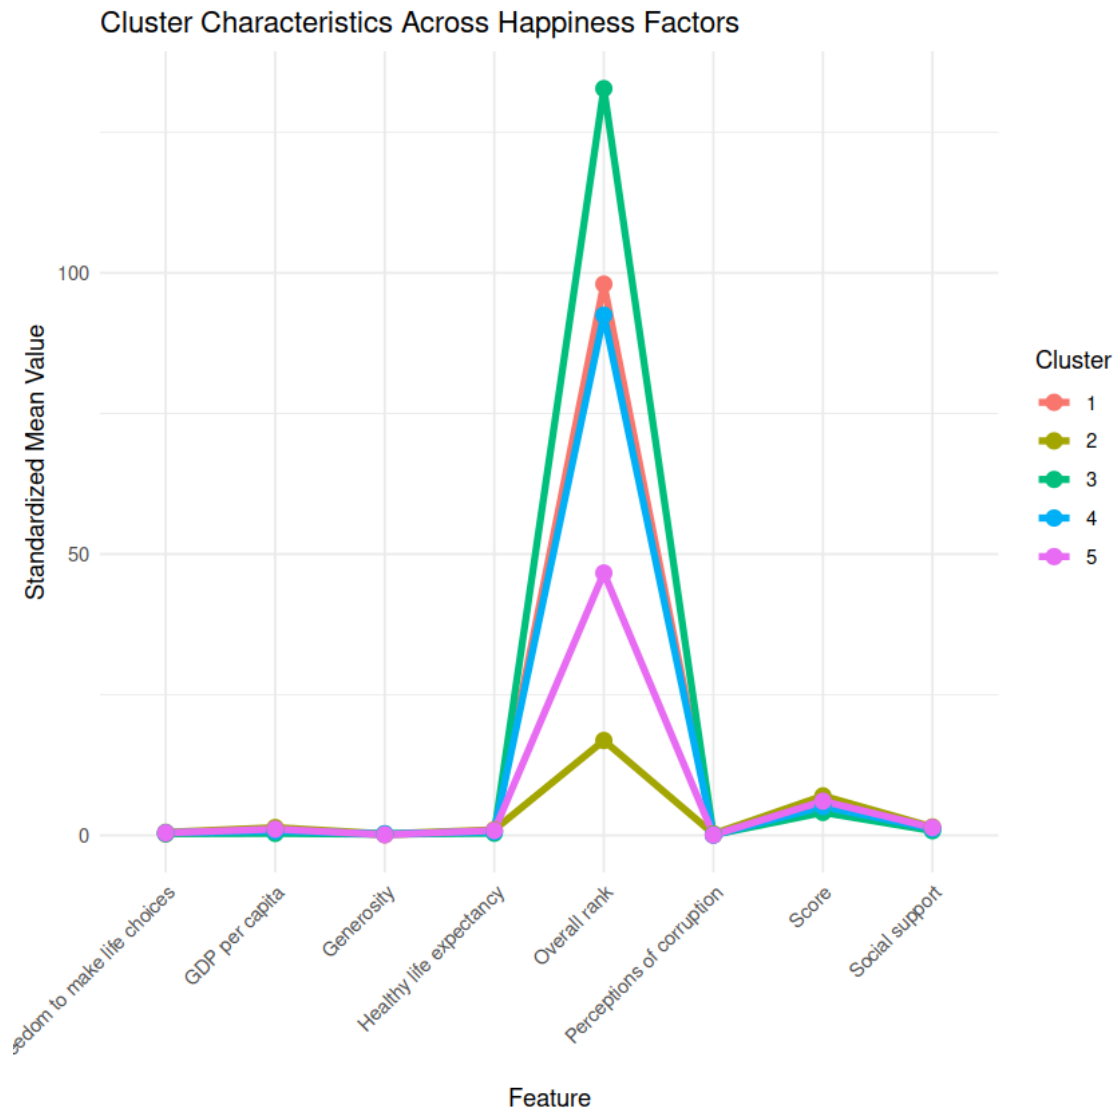
```
[27]: cluster_means_long <- cluster_means %>%
  pivot_longer(cols = -Cluster, names_to = "Feature", values_to = "Mean_Value")
```

```
[28]: ggplot(cluster_means_long, aes(x = Feature, y = Mean_Value, group = Cluster,
  ↪color = Cluster)) +
  geom_line(size = 1.5) +
  geom_point(size = 3) +
  labs(title = "Cluster Characteristics Across Happiness Factors",
```

```
y = "Standardized Mean Value") +
theme_minimal() +
theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Warning message:

"Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.  
Please use `linewidth` instead."



```
[29]: # Calculate and display cluster characteristics
cluster_profiles <- happiness %>%
  group_by(Cluster) %>%
  summarise(
    Count = n(),
```

17

```

    Avg_GDP = mean(`GDP per capita`, na.rm = TRUE),
    Avg_Social = mean(`Social support`, na.rm = TRUE),
    Avg_Health = mean(`Healthy life expectancy`, na.rm = TRUE),
    Avg_Freedom = mean(`Freedom to make life choices`, na.rm = TRUE),
    Avg_Generosity = mean(Generosity, na.rm = TRUE),
    Avg_Corruption = mean(`Perceptions of corruption`, na.rm = TRUE)
  )

```

```

[30]: # Print cluster profiles
      print(cluster_profiles)

```

```

# A tibble: 5 × 8
  Cluster Count Avg_GDP Avg_Social Avg_Health Avg_Freedom Avg_Generosity
  <fct>   <int>   <dbl>
1 1      31    0.968    1.20    0.759    0.259
0.106
2 2      23    1.40    1.49    0.990    0.548
0.275
3 3      36    0.351    0.824    0.387    0.301
0.203
4 4      21    0.771    1.17    0.661    0.457
0.290
5 5      45    1.12    1.39    0.867    0.449
0.130
# 1 more variable: Avg_Corruption <dbl>

```

```

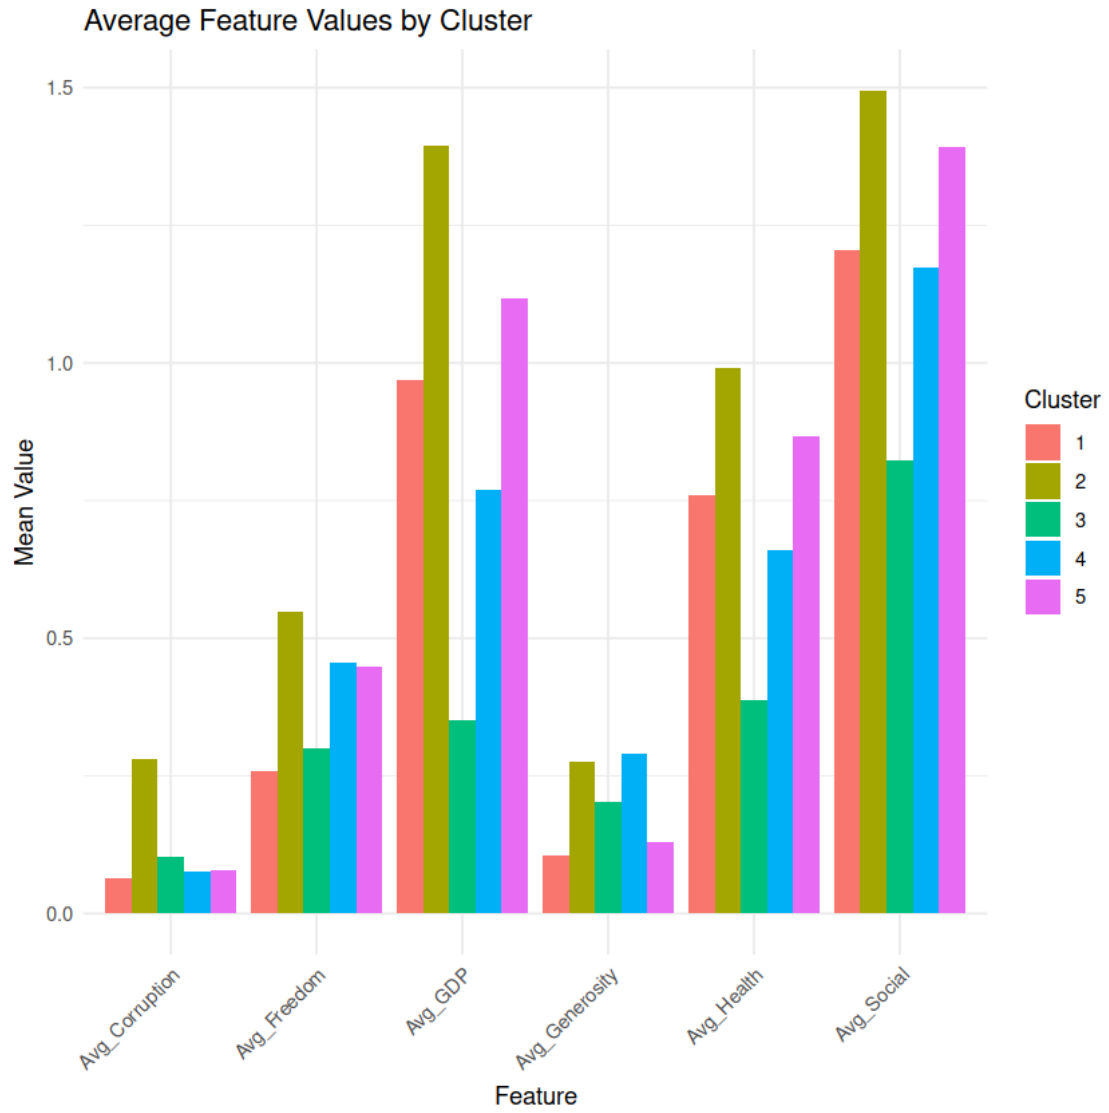
[31]: cluster_profiles_long <- cluster_profiles %>%
      select(-Count) %>%
      pivot_longer(cols = -Cluster, names_to = "Feature", values_to = "Mean_Value")

```

```

[32]: ggplot(cluster_profiles_long, aes(x = Feature, y = Mean_Value, fill = as.
      ↪factor(Cluster))) +
      geom_col(position = "dodge") +
      labs(title = "Average Feature Values by Cluster",
           y = "Mean Value",
           fill = "Cluster") +
      theme_minimal() +
      theme(axis.text.x = element_text(angle = 45, hjust = 1))

```



## Excercise3

April 21, 2025

```
[1]: library(tidyverse)
      library(cluster)
      library(gridExtra)
```

Attaching core tidyverse packages

```
tidyverse 2.0.0
  dplyr      1.1.4    readr      2.1.5
  forcats    1.0.0    stringr   1.5.1
  ggplot2    3.5.2    tibble    3.2.1
  lubridate  1.9.4    tidyr     1.3.1
  purrr      1.0.4
```

Conflicts

```
tidyverse_conflicts()
  dplyr::filter() masks stats::filter()
  dplyr::lag()    masks stats::lag()
Use the conflicted package
(<http://conflicted.r-lib.org/>) to force all conflicts to
become errors
```

Attaching package: 'gridExtra'

The following object is masked from 'package:dplyr':

combine

```
[2]: set.seed(1)
```

```
[3]: happiness <- read_csv("/home/asus/content/Notes/Semester 4/FDN Lab/Experiments/
  ↪Experiment 6/archive(11)/2019.csv")
```

Rows: 156 Columns: 9

## Column specification

Delimiter: ","

chr (1): Country

dbl (8): Overall rank, Score, GDP per capita, Social support, Healthy life e...

Use ``spec()`` to retrieve the full column specification for this data.

Specify the column types or set ``show_col_types = FALSE`` to quiet this message.

```
[4]: numeric_data <- happiness %>%
  select(
    "Overall rank", "Score", "GDP per capita", "Social support", "Healthy life_
    ↪expectancy", "Freedom to make life choices", "Generosity", "Perceptions of_
    ↪corruption"
  ) %>%
  scale()

[5]: rownames(numeric_data) <- happiness$`Country`

[6]: wcss <- map_dbl(1:10, ~ kmeans(numeric_data, ., nstart = 25)$tot.withinss)

[7]: avg_sil <- map_dbl(2:10, ~ {
  km <- kmeans(numeric_data, ., nstart = 25)
  silhouette_score <- silhouette(km$cluster, dist(numeric_data))
  mean(silhouette_score[, 3])
})

[8]: elbow_plot <- ggplot(data.frame(K = 1:10, WCSS = wcss), aes(K, WCSS)) +
  geom_line(color = "steelblue", size = 1.2) +
  geom_point(color = "red", size = 3) +
  labs(title = "Elbow Method (Optimal K)", x = "Number of Clusters (K)", y =_
  ↪"WCSS") +
  theme_minimal()

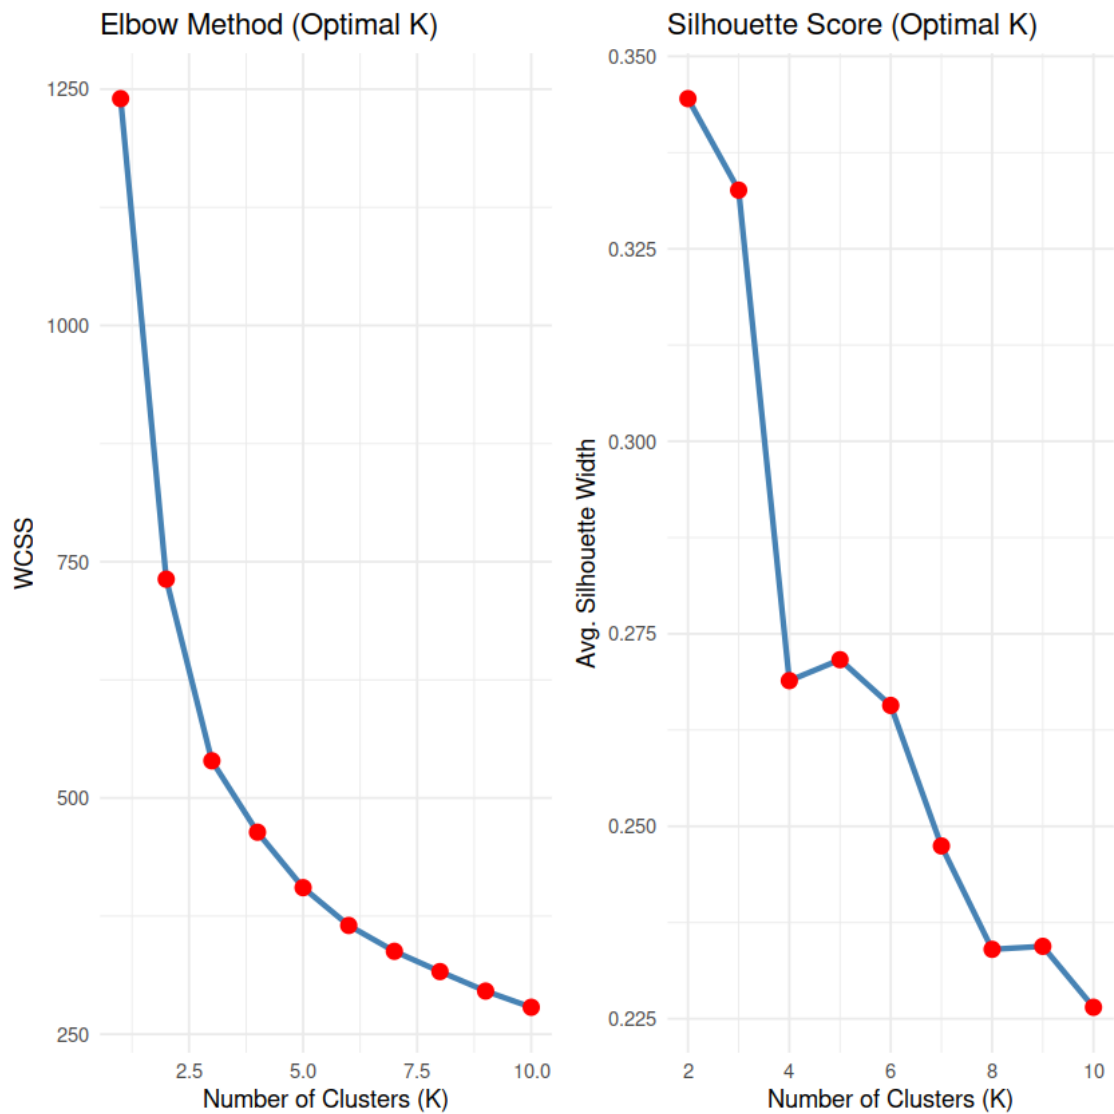
silhouette_plot <- ggplot(data.frame(K = 2:10, Silhouette = avg_sil), aes(K,_
  ↪Silhouette)) +
  geom_line(color = "steelblue", size = 1.2) +
  geom_point(color = "red", size = 3) +
  labs(title = "Silhouette Score (Optimal K)", x = "Number of Clusters (K)", y_
  ↪= "Avg. Silhouette Width") +
  theme_minimal()
```

```
grid.arrange(elbow_plot, silhouette_plot, ncol = 2)
```

Warning message:

"Using `size` aesthetic for lines was deprecated in ggplot2 3.4.0.

Please use `linewidth` instead."



```
[9]: k2 <- kmeans(numeric_data, centers = 2, nstart = 25)
      k3 <- kmeans(numeric_data, centers = 3, nstart = 25)
      k4 <- kmeans(numeric_data, centers = 4, nstart = 25)
      k5 <- kmeans(numeric_data, centers = 5, nstart = 25)
```

```
[10]: happiness$Cluster_K2 <- as.factor(k2$cluster)
      happiness$Cluster_K3 <- as.factor(k3$cluster)
```

```
happiness$Cluster_K4 <- as.factor(k4$cluster)
happiness$Cluster_K5 <- as.factor(k5$cluster)
```

```
[11]: plot_cluster_means <- function(km_result, title) {
  centers <- as.data.frame(km_result$centers)
  centers$Cluster <- factor(rownames(centers))

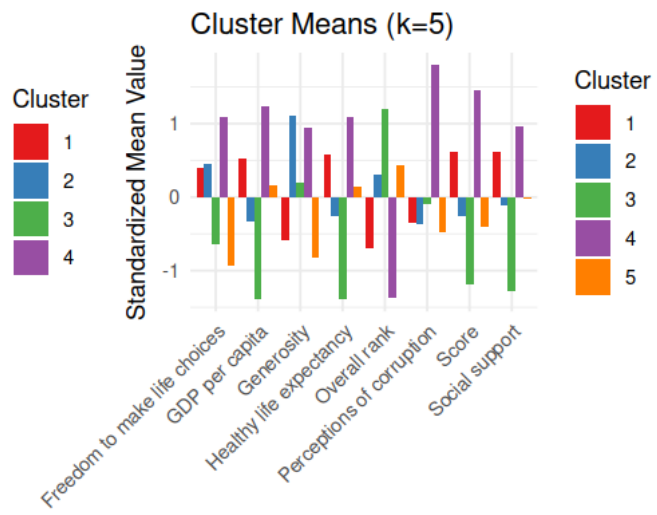

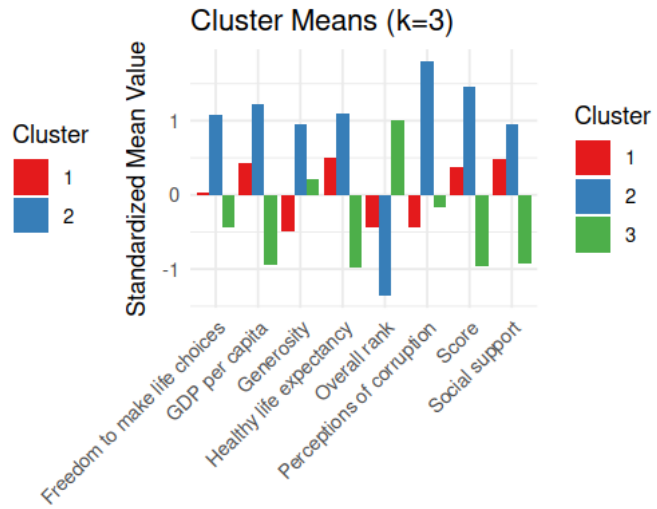

  centers_long <- centers %>%
    pivot_longer(cols = -Cluster, names_to = "Feature", values_to = "Mean_Value")

  ggplot(centers_long, aes(x = Feature, y = Mean_Value, fill = Cluster)) +
    geom_bar(stat = "identity", position = "dodge") +
    labs(title = title, y = "Standardized Mean Value", x = "") +
    theme_minimal() +
    theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
    scale_fill_brewer(palette = "Set1")
}
```

```
[12]: p2 <- plot_cluster_means(k2, "Cluster Means (k=2)")
p3 <- plot_cluster_means(k3, "Cluster Means (k=3)")
p4 <- plot_cluster_means(k4, "Cluster Means (k=4)")
p5 <- plot_cluster_means(k5, "Cluster Means (k=5)")

grid.arrange(p2, p3, p4, p5, ncol = 2)
```





# Experiment 7

April 21, 2025

```
[ ]: # Load required libraries
set.seed(1)
library(tidyverse)
library(caret)
library(glmnet)
library(mlbench)
library(randomForest)

[ ]: # Load Pima Indians Diabetes dataset
data("PimaIndiansDiabetes2")
df <- PimaIndiansDiabetes2

[ ]: # Check structure & missing values
glimpse(df)
summary(df)
df <- na.omit(df)

[ ]: preProc <- preProcess(df[, -9], method = c("center", "scale"))
df_scaled <- predict(preProc, df)

[ ]: df_scaled <- df_scaled %>% mutate(bmi_age_ratio = mass / age)

[ ]: cor_matrix <- cor(df_scaled %>% select(-diabetes))
corrplot(cor_matrix, method = "color", type = "upper", tl.cex = 0.7)

[ ]: ctrl <- rfeControl(functions = rfFuncs, method = "cv", number = 10)
rfe_result <- rfe(
  x = df_scaled %>% select(-diabetes),
  y = df_scaled$diabetes,
  sizes = 1:8, # Test subsets of 1 to 8 features
  rfeControl = ctrl
)

# Top selected features
print(rfe_result)
plot(rfe_result, type = c("g", "o"))
```

```
[ ]: x <- model.matrix(diabetes ~ ., df_scaled)[, -1] # Exclude intercept
y <- ifelse(df_scaled$diabetes == "pos", 1, 0)

# Fit LASSO
cv_lasso <- cv.glmnet(x, y, alpha = 1, family = "binomial")
plot(cv_lasso)

# Coefficients at optimal lambda
coef(cv_lasso, s = "lambda.min")
```

```
[ ]: trainIndex <- createDataPartition(df_scaled$diabetes, p = 0.8, list = FALSE)
train <- df_scaled[trainIndex, ]
test <- df_scaled[-trainIndex, ]
```

```
[ ]: model_all <- train(
  diabetes ~ .,
  data = train,
  method = "glm",
  family = "binomial",
  trControl = trainControl(method = "cv", number = 10)
)

pred_all <- predict(model_all, test)
confusionMatrix(pred_all, test$diabetes)
```

```
[ ]: model_selected <- train(
  diabetes ~ glucose + mass + bmi_age_ratio,
  data = train,
  method = "glm",
  family = "binomial",
  trControl = trainControl(method = "cv", number = 10)
)

# Predictions
pred_selected <- predict(model_selected, test)
confusionMatrix(pred_selected, test$diabetes)
```

```
[ ]:
```

# Experiment 8

April 21, 2025

```
[1]: library(tidyverse)
library(dplyr)
```

Attaching core tidyverse packages

```
tidyverse 2.0.0
dplyr      1.1.4      readr      2.1.5
forcats    1.0.0      stringr    1.5.1
ggplot2    3.5.2      tibble     3.2.1
lubridate  1.9.4      tidyr      1.3.1
purrr      1.0.4
Conflicts
```

```
tidyverse_conflicts()
dplyr::filter() masks stats::filter()
dplyr::lag()     masks stats::lag()
Use the conflicted package
(<http://conflicted.r-lib.org/>) to force all conflicts to
become errors
```

```
[2]: # Basic way to load a CSV
train <- read.csv("/home/asus/content/Notes/Semester 4/FDN Lab/Experiments/
↳Experiment 8/titanic/train.csv")

test <- read.csv("/home/asus/content/Notes/Semester 4/FDN Lab/Experiments/
↳Experiment 8/titanic/test.csv")
```

```
[3]: # Removing Unessacry Cols
train <- train %>% select(-one_of("Cabin", "Ticket", "Name", "Embarked"))
test <- test %>% select(-one_of("Cabin", "Ticket", "Name", "Embarked"))
```

```
[4]: train <- train %>% fill(everything(), .direction = "down")
test <- test %>% fill(everything(), .direction = "down")
```

```
[5]: X_train <- train %>% select(-Survived)
Y_train <- train %>% select(Survived)
```

```

[6]: train_df <- as_tibble(train) %>%
      mutate(Survived = train)

[7]: # Train the model
      logit_model <- glm(Survived ~ .,
                        data = train,
                        family = binomial)

[8]: predictions <- predict(logit_model, newdata = train, type = "response")

[9]: predicted_classes <- ifelse(predictions > 0.5, 1, 0)

[10]: ground_truth <- train$Survived

[11]: conf_matrix <- table(Predicted = predicted_classes, Actual = ground_truth)

[12]: print(conf_matrix)
# Actual
# Predicted    0    1
#              0 472 110
#              1   77 232

          Actual
Predicted    0    1
          0 472 110
          1   77 232

[13]: accuracy <- sum(diag(conf_matrix))/sum(conf_matrix)
      precision <- conf_matrix[2,2]/sum(conf_matrix[2,])
      recall <- conf_matrix[2,2]/sum(conf_matrix[,2])
      f1_score <- 2 * (precision * recall) / (precision + recall)

[14]: summary(logit_model)
      print(conf_matrix)
      cat("\nAccuracy:", round(accuracy, 3))
      cat("\nPrecision:", round(precision, 3))
      cat("\nRecall/Sensitivity:", round(recall, 3))
      cat("\nF1 Score:", round(f1_score, 3))
      cat("\nSpecificity:", round(conf_matrix[1,1]/sum(conf_matrix[,1]), 3))

```

Call:

```
glm(formula = Survived ~ ., family = binomial, data = train)
```

Deviance Residuals:

| Min     | 1Q      | Median  | 3Q     | Max    |
|---------|---------|---------|--------|--------|
| -2.6513 | -0.6196 | -0.4077 | 0.6269 | 2.6737 |

28

Coefficients:

|             | Estimate   | Std. Error | z value | Pr(> z )     |
|-------------|------------|------------|---------|--------------|
| (Intercept) | 4.6705255  | 0.5301556  | 8.810   | < 2e-16 ***  |
| PassengerId | 0.0000850  | 0.0003468  | 0.245   | 0.80639      |
| Pclass      | -1.0457412 | 0.1374434  | -7.609  | 2.77e-14 *** |
| Sexmale     | -2.8025118 | 0.2007943  | -13.957 | < 2e-16 ***  |
| Age         | -0.0338376 | 0.0067970  | -4.978  | 6.42e-07 *** |
| SibSp       | -0.3422950 | 0.1094887  | -3.126  | 0.00177 **   |
| Parch       | -0.1195347 | 0.1171594  | -1.020  | 0.30760      |
| Fare        | 0.0031898  | 0.0023918  | 1.334   | 0.18233      |

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 1186.66 on 890 degrees of freedom  
Residual deviance: 790.18 on 883 degrees of freedom  
AIC: 806.18

Number of Fisher Scoring iterations: 5

|           | Actual |     |
|-----------|--------|-----|
| Predicted | 0      | 1   |
| 0         | 472    | 110 |
| 1         | 77     | 232 |

Accuracy: 0.79  
Precision: 0.751  
Recall/Sensitivity: 0.678  
F1 Score: 0.713  
Specificity: 0.86